

# Parallel & Concurrent Haskell 2: The Par Monad

Simon Marlow

(Microsoft Research, Cambridge, UK)

# Recap

---

# Recap

---

- In part 1 we learned about
  - The Eval monad
    - nice simple primitives for introducing deterministic parallelism
    - minimal control over the evaluation order
  - Strategies
    - Adding parallelism over (lazy) data structures
    - Composability: combine Strategies into larger ones
    - Modularity: (e `using` s) separates the parallelism from the algorithm

# But...

---

# But...

---

- Lazy evaluation is the magic ingredient that bestows *modularity*, and thus forms the basis of Strategies.
  - but it can be tricky to deal with.

# But...

---

- Lazy evaluation is the magic ingredient that bestows *modularity*, and thus forms the basis of Strategies.
  - but it can be tricky to deal with.
- To use Strategies effectively, you need to understand things like
  - evaluation order (because the argument to `rpar` must be a lazy computation)
  - garbage collection (because the result of `rpar` must not be discarded)
  - In a sense this is all tricky *by design* because the Haskell language definition places no requirements on evaluation order or memory behaviour. Compilers are free to do what they like.

# But...

---

- Lazy evaluation is the magic ingredient that bestows *modularity*, and thus forms the basis of Strategies.
  - but it can be tricky to deal with.
- To use Strategies effectively, you need to understand things like
  - evaluation order (because the argument to `rpar` must be a lazy computation)
  - garbage collection (because the result of `rpar` must not be discarded)
  - In a sense this is all tricky *by design* because the Haskell language definition places no requirements on evaluation order or memory behaviour. Compilers are free to do what they like.
- Diagnosing performance problems can be hard

# The Par Monad

---

- Aim for a more direct programming model:
  - sacrifice “modularity via laziness”
  - Avoid the programmer having to think about when things are evaluated
    - ... hence avoid many common pitfalls
  - Modularity via *higher-order skeletons*
    - no laziness magic here, just higher-order functions and polymorphism
  - It’s a library written entirely in Haskell
    - Pure API outside, unsafePerformIO + forkIO inside
    - Write your own scheduler!




# The basic idea

---

- Think about your computation as a dataflow graph.

# Par expresses dynamic dataflow

---



# The **Par** Monad

Par is a monad for  
parallel computation

```
data Par  
instance Monad Par
```

```
runPar :: Par a -> a
```

```
fork :: Par () -> Par ()
```

```
data IVar
```

```
new :: Par (IVar a)
```

```
get :: IVar a -> Par a
```

```
put :: NFData a => IVar a -> a -> Par ()
```

# The **Par** Monad

Par is a monad for parallel computation

```
data Par
instance Monad Par
```

```
runPar :: Par a -> a
```

Parallel computations are pure (and hence deterministic)

```
fork :: Par () -> Par ()
```

```
data IVar
```

```
new :: Par (IVar a)
```

```
get :: IVar a -> Par a
```

```
put :: NFData a => IVar a -> a -> Par ()
```

# The **Par** Monad

```
data Par
instance Monad Par
```

Par is a monad for  
parallel computation

```
runPar :: Par a -> a
```

Parallel computations  
are pure (and hence  
deterministic)

```
fork :: Par () -> Par ()
```

forking is *explicit*

```
data IVar
new :: Par (IVar a)
get :: IVar a -> Par a
put :: NFData a => IVar a -> a -> Par ()
```

# The **Par** Monad

```
data Par  
instance Monad Par
```

Par is a monad for  
parallel computation

```
runPar :: Par a -> a
```

Parallel computations  
are pure (and hence  
deterministic)

```
fork :: Par () -> Par ()
```

forking is *explicit*

```
data IVar
```

results are communicated  
through IVars

```
new :: Par (IVar a)
```

```
get :: IVar a -> Par a
```

```
put :: NFData a => IVar a -> a -> Par ()
```

# A couple of things to bear in mind

---

# A couple of things to bear in mind

---

- *put is fully strict*

```
put :: NFData a => IVar a -> a -> Par ()
```

- all values communicated through **IVars** are fully evaluated
  - The programmer can tell where the computation is happening, and hence reason about the parallelism
- (there is a head-strict version **put\_** but we won't be using it)



# A couple of things to bear in mind

---

- *put is fully strict*


```
put :: NFData a => IVar a -> a -> Par ()
```

- all values communicated through *IVars* are fully evaluated
  - The programmer can tell where the computation is happening, and hence reason about the parallelism
- (there is a head-strict version *put\_* but we won't be using it)
- *put twice on the same IVar is an error*
  - This is a requirement for *Par* to be deterministic

# How does this make a dataflow graph?

```
do v <- new  
  fork $ put v (f x)  
  get v
```


- fork creates a new node in the graph
- get creates a new edge
  - from the node containing the put
  - to the node containing the get



# A bit more complex...

```
do v1 <- new
   v2 <- new
   fork $ put v1 (f x)
   fork $ put v2 (g x)
   get v1
   get v2
   return (v1 + v2)
```

- **runPar** evaluates the graph
- nodes with no dependencies between them can execute in parallel




# A bit more complex...

```
do v1 <- new
   v2 <- new
   fork $ put v1 (f x)
   fork $ put v2 (g x)
   get v1
   get v2
   return (v1 + v2)
```

- **runPar** evaluates the graph
- nodes with no dependencies between them can execute in parallel

Parallel!



# Back to the Sudoku example

---

- The sequential code:

```
main :: IO ()
main = do
    [f] <- getArgs
    grids <- fmap lines $ readFile f
    print $ length $ filter isJust $ map solve grids
```

```
solve :: String -> Maybe Grid
```

# Sudoku solver, version 2

---

- Divide the work in two:

```
import Control.Monad.Par

main :: IO ()
main = do
    [f] <- getArgs
    grids <- fmap lines $ readFile f

    let (as,bs) = splitAt (length grids `div` 2) grids

    print $ length $ filter isJust $ runPar $ do
        i1 <- new
        i2 <- new
        fork $ put i1 (map solve as)
        fork $ put i2 (map solve bs)
        as' <- get i1
        bs' <- get i2
        return (as' ++ bs')
```

# Compile it for parallel execution

---

```
$ ghc --make -O2 sudoku-par2.hs -rtsopts -threaded
[1 of 2] Compiling Sudoku          ( Sudoku.hs, Sudoku.o )
[2 of 2] Compiling Main            ( sudoku-par2.hs, sudoku-par2.o )
Linking sudoku-par2 ...
$
```

# Run it on 2 processors

```
> ./sudoku-par2 sudoku17.1000.txt +RTS -s -N2
./sudoku-par2 sudoku17.1000.txt +RTS -s -N2
1000
  2,400,207,256 bytes allocated in the heap
  49,191,144 bytes copied during GC
  2,669,416 bytes maximum residency (7 sample(s))
  339,904 bytes maximum slop
    9 MB total memory in use (0 MB lost due to fragmentation)
```

...

INIT	time	0.00s	( 0.00s elapsed)
MUT	time	2.24s	( 1.79s elapsed)
GC	time	1.11s	( 0.20s elapsed)
EXIT	time	0.00s	( 0.00s elapsed)
Total	time	3.34s	( <b>1.99s elapsed</b> )

...

Speedup, yay!




File View Help



Key Traces Bookmarks

Timeline

- running
- GC
- create thread
- run spark
- thread runnable
- seq GC req
- par GC req
- migrate thread
- thread wakeup
- shutdown



Events

0.001139s startup: 2 capabilities  
0.001453s cap 1: creating thread 1  
0.001454s cap 1: thread 1 is runnable  
0.001457s cap 1: running thread 1  
0.001564s cap 1: stopping thread 1 (making a foreign call)  
0.001566s cap 1: running thread 1  
0.001573s cap 1: stopping thread 1 (making a foreign call)

# parMap

- Let's use the **Par** monad to define the **parMap** pattern. First expand our vocabulary a bit:


```
spawn :: Par a -> Par (IVar a)
spawn p = do r <- new
           fork $ p >>= put r
           return r
```

- now define **parMap**:

```
parMap :: NFData b => (a -> b) -> [a] -> Par [b]
parMap f as = do
  ibs <- mapM (spawn . return . f) as
  mapM get ibs
```

# What is the dataflow graph?

---



# Parallel sudoku solver version 3

---

```
main :: IO ()
main = do
  [f] <- getArgs
  grids <- fmap lines $ readFile f
  print $ length $ filter isJust $
    runPar $ parMap solve grids
```

- How does it perform?

# sudoku-par3 on 2 cores

```
./sudoku-par3 sudoku17.1000.txt +RTS -N2 -s
1000
  2,400,973,624 bytes allocated in the heap
  50,751,248 bytes copied during GC
  2,654,008 bytes maximum residency (6 sample(s))
  368,256 bytes maximum slop
    9 MB total memory in use (0 MB lost due to fragmentation)

...

INIT   time    0.00s   (  0.00s elapsed)
MUT    time    2.06s   (  1.47s elapsed)
GC      time    1.29s   (  0.21s elapsed)
EXIT   time    0.00s   (  0.00s elapsed)
Total  time    3.36s   (  1.68s elapsed)
```

- Speedup:  $3.02/1.68 = 1.79$


File View Help



Key Traces Bookmarks

- running
- GC
- create thread
- run spark
- thread runnable
- seq GC req
- par GC req
- migrate thread
- thread wakeup
- shutdown

Timeline



Events

1.691729s	cap 0: GC idle
1.691729s	cap 0: GC done
1.691749s	cap 1: finished GC
1.691763s	cap 0: running thread 2
1.691827s	cap 0: stopping thread 2 (thread finished)
1.691851s	cap 0: shutting down
1.691853s	cap 1: shutting down

# Granularity

---

- Granularity = size of the tasks
  - Too small, and the overhead of `fork/get/put` will outweigh the benefits of parallelism
  - Too large, and we risk underutilisation (see `sudoku-par2.hs`)
  - The range of “just right” is often quite wide
- Let’s test that. How do we change the granularity?

# parMap with variable granularity

```
parMapChunk :: NFData b => Int -> (a -> b) -> [a] -> Par [b]
parMapChunk n f xs = do
  xss <- parMap (map f) (chunk n xs)
  return (concat xss)

chunk :: Int -> [a] -> [[a]]
chunk _ [] = []
chunk n xs = as : chunk n bs
  where (as,bs) = splitAt n xs
```

- split the list into chunks of size  $n$
- Each node processes  $n$  elements
- (this isn't in the library, but it should be)



# Final version of sudoku: chunking

---

- sudoku-par4.hs

```
main :: IO ()
main = do
    [f,n] <- getArgs
    grids <- fmap lines $ readFile f
    print $ length $ filter isJust $
        runPar $ parMapChunk (read n) solve grids
```

# Results with sudoku17.16000.txt

No chunks (sudoku-par3):

Total time 43.71s ( 43.73s elapsed)

chunk 100, -N1:

Total time 44.43s ( 44.44s elapsed)

No chunks, -N8:

Total time 67.73s ( 8.38s elapsed)

(5.21x)

chunk 10, -N8:

Total time 61.62s ( 7.74s elapsed)

(5.64x)

chunk 100, -N8:

Total time 60.81s ( 7.73s elapsed)

(5.65x)

chunk 1000, -N8:

Total time 61.74s ( 7.88s elapsed)

(5.54x)

# Granularity: conclusion

---

- Use `parListChunk` if your tasks are too small
- If your tasks are too large, look for ways to divide the work and add more parallelism
- If the number of tasks is less than 10 times the number of cores, that is probably too few

# Enough about sudoku!

---

- We've been dealing with flat parallelism so far
- What about other common patterns, such as divide and conquer?

# Examples

---

- Divide and conquer parallelism:

```
parfib :: Int -> Int -> Par Int
parfib n
  | n <= 2    = return 1
  | otherwise = do
    x <- spawn $ parfib (n-1)
    y <- spawn $ parfib (n-2)
    x' <- get x
    y' <- get y
    return (x' + y')
```

# Note...

---

# Note...

---

- We have to thread the Par monad to all the sites we might want to spawn or fork.

# Note...

---

- We have to thread the Par monad to all the sites we might want to spawn or fork.
- Why? Couldn't we just call a new runPar each time?

```
runPar :: Par a -> a
```



# Note...

---

- We have to thread the Par monad to all the sites we might want to spawn or fork.
- Why? Couldn't we just call a new runPar each time?

```
runPar :: Par a -> a
```

- Each **runPar**:
  - Waits for all its subtasks to finish before returning (necessary for determinism)
  - Fires up a new gang of N threads and creates scheduling data structures: it's expensive
  - So we do want to thread the **Par** monad around


# Granularity in divide-and-conquer

- If you try to run this, performance will be terrible:

```
parfib :: Int -> Par Int
parfib n
  | n <= 2    = return 1
  | otherwise = do
    x <- spawn $ parfib (n-1)
    y <- spawn $ parfib (n-2)
    x' <- get x
    y' <- get y
    return (x' + y')
```

- For a start, it's 50x slower than the sequential version
  - overhead of the Par monad

- As we saw before, when our tasks are too small we need to increase the granularity
- Here there's no obvious place to do chunking
- Instead we want to set a threshold for task creation



- parfib takes an extra parameter, the threshold
- below the threshold, we use the sequential fib
- a threshold of e.g. 25 is enough to give almost perfect speedup

```
parfib :: Int -> Int -> Par Int
parfib n t
  | n <= 2      = return 1
  | n <= t      = fib n
  | otherwise = do
    x <- spawn $ parfib (n-1) t
    y <- spawn $ parfib (n-2) t
    x' <- get x
    y' <- get y
    return (x' + y')
```

```
fib :: Int -> Int
fib n = ...
```

- parfib isn't a very realistic example
- let's try sorting instead
- mergesort:
  - divide the list into two
  - sort each half
  - merge the results

```
parsort :: Int -> [Integer] -> Par [Integer]
parsort n [] = return []
parsort n [x] = return [x]
parsort 0 xs = return (sort xs)
parsort n xs = do
    let (as,bs) = split xs
    l <- spawn $ parsort (n-1) as
    r <- spawn $ parsort (n-1) bs
    ls <- get l
    rs <- get r
    return (merge ls rs)
```

- **n** is the threshold
- **sort** is the sequential sort

# Skeletons

---

- Parallelism often fits a well-known pattern
- We've seen two common patterns so far:
  - parallel map
  - divide-and-conquer
- The idea of a skeleton is to abstract the pattern as a reusable higher-order function
- **parMap** is already a skeleton

# Divide and conquer as a skeleton

```
divConq :: NFData sol
    => (prob -> Bool)           -- indivisible?
    -> (prob -> (prob,prob))    -- split into subproblems
    -> (sol -> sol -> sol)     -- join solutions
    -> (prob -> sol)           -- solve a subproblem
    -> (prob -> sol)
```

```
divConq indiv split join f prob
= runPar $ go prob
  where
    go prob
      | indiv prob = return (f prob)
      | otherwise = do
          let (a,b) = split prob
          i <- spawn $ go a
          j <- spawn $ go b
          a <- get i
          b <- get j
          return (join a b)
```

- Using the skeleton
- Our “prob” is `(Int,[Integer])`
  - i.e. pair the threshold counter with the list

```
parsort :: Int -> [Integer] -> [Integer]
parsort thresh xs
  = divConq indiv divide merge (sort . snd) (thresh,xs)
  where
    indiv (n,xs) = n == 0

    divide (n,xs) = ((n-1, as), (n-1, bs))
      where (as,bs) = split xs
```

- Nice compact definition of parallel sorting
- Important: the details of the parallelism are hidden in `divConq` (we could have used `Strategies`)



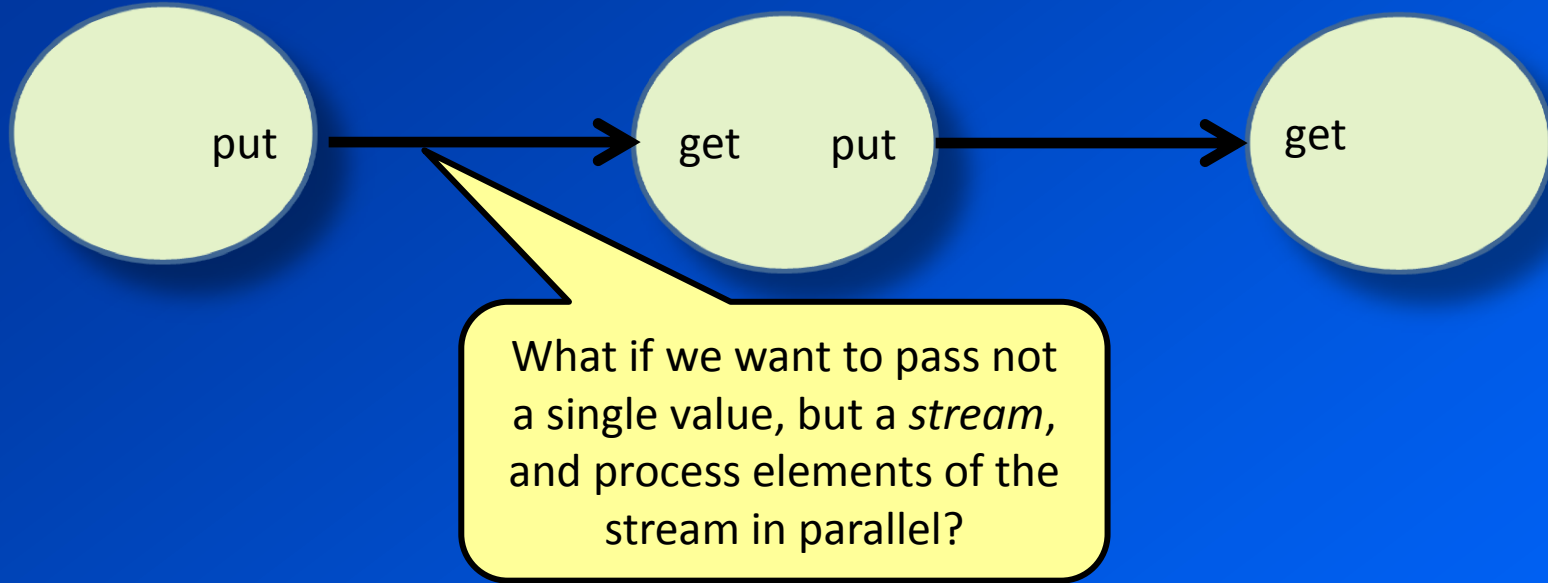
# Rule of thumb

---

- Try to separate the *application code* from the *parallel coordination* by using higher-order skeletons
- Good abstraction facilities lead to modularity

break...

# Pipeline parallelism



Or to put it another way:

- if we have a computation that produces a list
- and another one that consumes it
- how can we overlap these with the Par monad?

# IList and Stream

```
data IList a = Null
             | Cons { hd :: a
                     , tl :: Stream a }

type Stream a = IVar (IList a)
```

- Stream is a “lazy list” in the **Par** monad
  - but we’re being explicit about where the laziness is
- We need a way to:
  - Generate a new Stream
  - Process a stream (map, filter)
  - Consume a Stream (fold)
- Plugging these together gives us parallel pipeline processing.
- Stream code is in <code/euler35/Stream.hs>

# Generate a Stream

- One way: generate a stream from a real lazy list:

```
fromList :: NFData a => [a] -> Par (Stream a)
fromList xs =
    do var <- new
      fork $ loop xs var
    return var
where
    loop [] var = put var Null
    loop (x:xs) var =
        do tail <- new
          put var (Cons x tail)
          loop xs tail
```

Strict!

# Filter a Stream

```
streamFilter :: NFData a => (a -> Bool) -> Stream a
              -> Par (Stream a)
streamFilter p instr = do
  outstr <- new
  fork $ loop instr outstr
  return outstr
where
  loop instr outstr = do
    v <- get instr
    case v of
      Null -> put outstr Null
      Cons x instr'
        | p x -> do
          tail <- new
          put_ outstr (Cons x tail)
          loop instr' tail
        | otherwise -> do
          loop instr' outstr
```

# Consume a stream

---

- Analogue of **foldl**:

```
streamFold :: (a -> b -> a) -> a -> Stream b -> Par a
streamFold fn acc instrm =
  do ilst <- get instrm
  case ilst of
    Null      -> return acc
    Cons h t  -> streamFold fn (fn acc h) t
```

- This version is not strict – maybe it should be?

# Pipeline example

- Project Euler problem 35: “Find all the *circular* primes below 1,000,000”. A circular prime is one in which all rotations of its digits are also prime.

```
main :: IO ()
main = print $ runPar $ do
  s1 <- streamFromList (takeWhile (<1000000) primes)
  s2 <- streamFilter circular s1
  streamFold (\a _ -> a + 1) 0 s2
```

- Full code is in [code/euler35/euler35.hs](#)
- Achieves 1.85 speedup vs. the sequential version on 2 cores (does not scale further)
- Beware: this is not suitable for working with streams that do not fit in memory, since there is nothing preventing the producer from producing elements too fast



# Dataflow problems

---

- Par really shines when the problem is easily expressed as a dataflow graph, particularly an irregular or dynamic graph (e.g. shape depends on the program input)
- Identify the nodes and edges of the graph
  - each node is created by **fork**
  - each edge is an **lVar**


# Example

---

- Consider typechecking a functional program
- A set of bindings of the form  $x = e$
- To typecheck a binding:
  - input: the types of the variables mentioned in  $e$
  - output: the type of  $x$
- So this is a dataflow graph
  - a node represents the typechecking of a binding
  - the types of identifiers flow down the edges
  - It's a *dynamic* dataflow graph: we don't know the shape beforehand

# Example

```
f = ...  
g = ... f ...  
h = ... f ...  
j = ... g ... h ...
```



# Implementation

---

- We parallelised an existing type checker (nofib/infer).
- Algorithm works on a single term:

```
data Term = Let VarId Term Term | ...
```

- So we parallelise checking of the top-level Let bindings.

```
let x1 = e1 in  
let x2 = e2 in  
let x3 = e3 in  
...
```

# The parallel type inferencer

---

- Given:

```
inferTopRhs :: Env -> Term -> PolyType  
makeEnv    :: [(VarId,Type)] -> Env
```

- We need a type environment:

```
type TopEnv = Map VarId (IVar PolyType)
```

- The top-level inferencer has the following type:

```
inferTop :: TopEnv -> Term -> Par MonoType
```

# Parallel type inference

```
inferTop :: TopEnv -> Term -> Par MonoType
inferTop topenv (Let x u v) = do
  vu <- new

  fork $ do
    let fu = Set.toList (freeVars u)
    tfu <- mapM (get . fromJust . flip Map.lookup topenv) fu
    let aa = makeEnv (zip fu tfu)
    put vu (inferTopRhs aa u)

  inferTop (Map.insert x vu topenv) v

inferTop topenv t = do
  -- the boring case: invoke the normal sequential
  -- type inference engine
```


# Results

```
let id = \x.x in
  let x = \f.f id id in
  let x = \f . f x x in
  let x = \f . f x x in
  let x = \f . f x x in
  ...
  let x = let f = x in \z . z in
  let y = \f.f id id in
  let y = \f . f y y in
  let y = \f . f y y in
  let y = \f . f y y in
  ...
  let x = let f = y in \z . z in
  \f. let g = \a. a x y in f
```

- -N1: 1.12s
- -N2: 0.60s (1.87x speedup)
- available parallelism depends on the input: these bindings only have two branches

# K-Means

- A data-mining algorithm, to identify clusters in a data set.





# K-Means

---

- We use a heuristic technique (Lloyd's algorithm), based on iterative refinement.
  1. Input: an initial guess at each cluster location
  2. Assign each data point to the cluster to which it is closest
  3. Find the *centroid* of each cluster (the average of all points)
  4. repeat 2-3 until clusters stabilise
- Making the initial guess:
  1. Input: number of clusters to find
  2. Assign each data point to a random cluster
  3. Find the centroid of each cluster
- Careful: sometimes a cluster ends up with no points!

# K-Means: basics

```
data Vector = Vector Double Double
```

```
addVector :: Vector -> Vector -> Vector
```

```
addVector (Vector a b) (Vector c d) = Vector (a+c) (b+d)
```

```
data Cluster = Cluster
```

```
{
```

```
    clId      :: !Int,
```

```
    clCount   :: !Int,
```

```
    clSum     :: !Vector,
```

```
    clCent    :: !Vector
```

```
}
```

```
sqDistance :: Vector -> Vector -> Double
```

```
-- square of distance between vectors
```

```
makeCluster :: Int -> [Vector] -> Cluster
```

```
-- builds Cluster from a set of points
```

# K-Means:

```
assign
  :: Int          -- number of clusters
  -> [Cluster]    -- clusters
  -> [Vector]     -- points
  -> Array Int [Vector] -- points assigned to clusters

makeNewClusters :: Array Int [Vector] -> [Cluster]
  -- takes result of assign, produces new clusters

step :: Int -> [Cluster] -> [Vector] -> [Cluster]
step nclusters clusters points =
  makeNewClusters (assign nclusters clusters points)
```

- **assign** is step 2 (assign points to clusters)
- **makeNewClusters** is step 3 (compute average of points to get new clusters)
- **step** is (2,3) – one iteration

# Putting it together.. sequentially

```
kmeans_seq :: Int -> [Vector] -> [Cluster] -> IO [Cluster]
kmeans_seq nclusters points clusters = do
  let
    loop :: Int -> [Cluster] -> IO [Cluster]
    loop n clusters | n > tooMany = return clusters
    loop n clusters = do
      hPrintf stderr "iteration %d\n" n
      hPutStr stderr (unlines (map show clusters))
      let clusters' = step nclusters clusters points
      if clusters' == clusters
      then return clusters
      else loop (n+1) clusters'
  --
  loop 0 clusters
```

# Parallelise makeNewClusters?

---

```
makeNewClusters :: Array Int [Vector] -> [Cluster]
makeNewClusters arr =
  filter ((>0) . clCount) $
    [ makeCluster i ps | (i,ps) <- assocs arr ]
```

- essentially a map over the clusters
- number of clusters is small
- not enough parallelism here – grains are too large, fan-out is too small

# How to parallelise?

- Parallelise assign?

```
assign :: Int -> [Cluster] -> [Vector] -> Array Int [Vector]
assign nclusters clusters points =
    accumArray (flip (:)) [] (0, nclusters-1)
        [ (clId (nearest p), p) | p <- points ]
    where
        nearest p = ...
```

- essentially map/reduce: map nearest + **accumArray** to attach each point to its cluster
- the **map** parallelises, but **accumArray** doesn't
- could divide into chunks... but is there a better way?

# Sub-divide the data

---

- Suppose we divided the data set in two, and called **step** on each half
- We would need a way to combine the results:

```
step n cs (as ++ bs) == step n cs as `combine` step n cs bs
```

- but what is **combine**?

```
combine :: [Cluster] -> [Cluster] -> [Cluster]
```

- assuming we can match up cluster pairs, we just need a way to combine two clusters

# Combining clusters

- A cluster is notionally a set of points
- Its *centroid* is the average of the points
- A Cluster is represented by its centroid:

```
data Cluster = Cluster
    {
        clId      :: !Int,
        clCount   :: !Int,      -- num of points
        clSum      :: !Vector,  -- sum of points
        clCent     :: !Vector   -- clSum / clCount
    }
```

- but note that we cached **clCount** and **clSum**
- these let us merge two clusters and recompute the centroid in  $O(1)$



# Combining clusters

---

- So using

```
combineClusters :: Cluster -> Cluster -> Cluster
```

- we can define

```
reduce :: Int -> [[Cluster]] -> [Cluster]
```

- (see notes for the code; straightforward)
- now we can express K-Means as a map/reduce

# Final parallel implementation

```
kmeans_par :: Int -> Int -> [Vector] -> [Cluster] -> IO [Cluster]
kmeans_par chunks nclusters points clusters = do
  let chunks = split chunks points
  let
    loop :: Int -> [Cluster] -> IO [Cluster]
    loop n clusters | n > tooMany = return clusters
    loop n clusters = do
      hPrintf stderr "iteration %d\n" n
      hPutStr stderr (unlines (map show clusters))
      let
        new_clusterss =
          runpar $ parMap (step nclusters clusters) chunks

        clusters' = reduce nclusters new_clusterss


      if clusters' == clusters
      then return clusters
      else loop (n+1) clusters'
  --
  loop 0 clusters
```

# What chunk size?

---

- Divide data by number of processors?
  - No! Static partitioning could lead to poor utilisation (see earlier)
  - there's no need to have such large chunks, the RTS will schedule smaller work items across the available cores

- Results for 170000 2-D points, 4 clusters, 1000 chunks



# Further thoughts

---

- We had to restructure the algorithm to make the maximum amount of parallelism available
  - map/reduce
  - move the branching point to the top
  - make reduce as cheap as possible
  - a tree of reducers is also possible
- Note that the parallel algorithm is data-local – this makes it particularly suitable for distributed parallelism (indeed K-Means is commonly used as an example of distributed parallelism).
- But be careful of static partitioning

# Thoughts to take away...

---

# Thoughts to take away...

---

- Making your program faster is the goal
  - Parallelism is just one way to achieve that
    - it might not be the easiest way!
  - However, designing your code with parallelism in mind should ensure that it can ride Moore's law a bit longer
  - good:
    - maps and trees
  - suspicious:
    - folds (but associative folds are OK)
    - lists (operations on lists themselves are serial, but operations on the elements of a list can be parallelised)

# Exercises

---

- Don't use the printout!
- <http://community.haskell.org/~simonmar/lab-exercises-cadarache.pdf>
- includes instructions for downloading the sample code
- Lab 1 covers the first two lectures
- Enjoy!