

Parallel & Concurrent Haskell 3: Concurrent Haskell

Simon Marlow

Concurrent Haskell

- Recap:
 - concurrent programming is about *threads of control*
 - concurrency is necessary for dealing with multiple sources of input/output:
 - network connections
 - GUI, user input
 - database connections
 - *threads of control* let you handle multiple input/output sources in a *modular* way: the code for each source is written separately

Summary

- In this part of the course we're going to cover
 - Basic concurrency
 - forking threads
 - communication and synchronisation
 - Asynchronous exceptions
 - cancellation
 - timeouts

Forking threads

```
forkIO :: IO () -> IO ThreadId
```

- creates a new thread to run the IO ()
- new thread runs “at the same time” as the current thread and other threads

Interleaving example

```
import Control.Concurrent
import Control.Monad
import System.IO

main = do
  hSetBuffering stdout NoBuffering
  forkIO (forever (putChar 'A'))
  forkIO (forever (putChar 'B'))
  threadDelay (10^6)
```

```
forkIO :: IO () -> ThreadId

forever :: m a -> m a
putChar :: Char -> IO ()
threadDelay :: Int -> IO ()
```

```
$ ghc fork.hs
[1 of 1] Compiling Main                ( fork.hs, fork.o )
Linking fork ...
$ ./fork | tail -c 300
AAAAAAAAABABABABABABABABABABABABABABABABABABABABABABABAB
ABABABABABABABABABABABABABABABABABABABABABABABABABABAB
ABABABABABABABABABABABABABABABABABABABABABABABABABABAB
ABABABABABABABABABABABABABABABABABABABABABABABABABABAB
$
```

ThreadId

```
forkIO :: IO () -> IO ThreadId
```

– what can you do with a ThreadId?

- check status with `GHC.Conc.threadStatus` (useful for debugging):

```
> import Control.Concurrent
> do { t <- forkIO (threadDelay (10^6)); GHC.Conc.threadStatus t }
ThreadIdRunning
> do { t <- forkIO (threadDelay (10^6)); yield; GHC.Conc.threadStatus t }
ThreadIdBlocked BlockedOnMVar
```

– Also:

- compare for equality
- kill / send exceptions (later...)

A note about performance

- GHC's threads are *lightweight*

```
> ./Main 1000000 1 +RTS -s
```

```
Creating pipeline with 1000000 processes in it.
```

```
Pumping a single message through the pipeline.
```

```
Pumping a 1 messages through the pipeline.
```

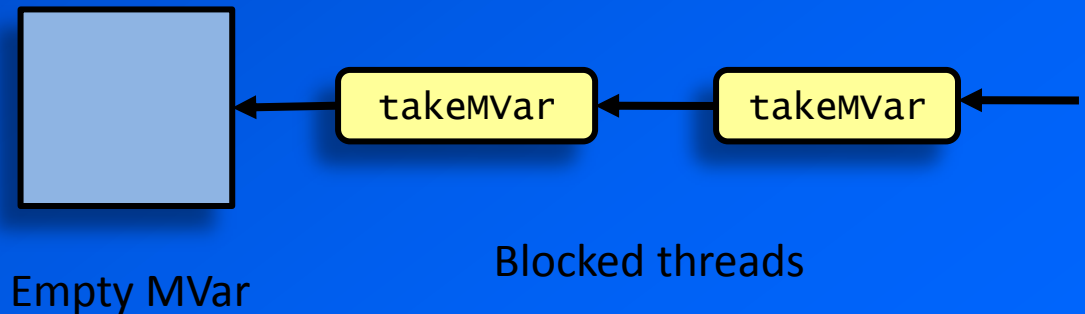
n	create	pump1	pump2	create/n	pump1/n	pump2/n
	s	s	s	us	us	us
1000000	5.980	1.770	1.770	5.98	1.77	1.77

- 10^6 threads requires 1.5Gb – 1.5k/thread
 - most of that is stack space, which grows/shrinks on demand
- cheap threads makes it feasible to use them liberally, e.g. one thread per client in a server

Communication: MVars

- MVar is the basic communication primitive in Haskell.


```
data MVar a -- abstract  
  
newEmptyMVar :: IO (MVar a)  
takeMVar     :: MVar a -> IO a  
putMVar      :: MVar a -> a -> IO ()
```



Communication: MVars

- MVar is the basic communication primitive in Haskell.


```
data MVar a    -- abstract  
  
newEmptyMVar   :: IO (MVar a)  
takeMVar       :: MVar a -> IO a  
putMVar        :: MVar a -> a -> IO ()
```



Communication: MVars

- MVar is the basic communication primitive in Haskell.

```
data MVar a -- abstract  
  
newEmptyMVar :: IO (MVar a)  
takeMVar     :: MVar a -> IO a  
putMVar      :: MVar a -> a -> IO ()
```



Communication: MVars

- MVar is the basic communication primitive in Haskell.

```
data MVar a -- abstract  
  
newEmptyMVar :: IO (MVar a)  
takeMVar     :: MVar a -> IO a  
putMVar      :: MVar a -> a -> IO ()
```

putMVar x

x <- takeMVar



Empty MVar

← takeMVar ←


Blocked threads

Communication: MVars

- MVar is the basic communication primitive in Haskell.

```
data MVar a -- abstract

newEmptyMVar :: IO (MVar a)
takeMVar     :: MVar a -> IO a
putMVar      :: MVar a -> a -> IO ()
```




Communication: MVars

- MVar is the basic communication primitive in Haskell.

```
data MVar a -- abstract

newEmptyMVar :: IO (MVar a)
takeMVar     :: MVar a -> IO a
putMVar      :: MVar a -> a -> IO ()
```



- And conversely: **putMVar** blocks when the MVar is full.

Examples

```
do m <- newEmptyMVar  
  forkIO $ putMVar m 'a'  
  takeMVar m
```

Examples

```
do m <- newEmptyMVar  
  forkIO $ putMVar m 'a'  
  takeMVar m
```

```
> do m <- newEmptyMVar; forkIO $ putMVar m 'a'; takeMVar m  
'a'  
>
```

Examples

```
do m <- newEmptyMVar  
  forkIO (takeMVar m >>= print)  
  putMVar m 'a'
```


Examples

```
do m <- newEmptyMVar  
  forkIO (takeMVar m >=> print)  
  putMVar m 'a'
```

```
> do m <- newEmptyMVar; forkIO (takeMVar m >=> print);  
  putMVar m 'a'  
'a'>
```

Examples

```
do m <- newEmptyMVar  
  forkIO (takeMVar m >>= print)  
  putMVar m 'a'
```

```
> do m <- newEmptyMVar; forkIO (takeMVar m >>= print);  
  putMVar m 'a'  
'a'>
```

```
> do m <- newEmptyMVar; forkIO (threadDelay 1000 >>  
  takeMVar m >>= print); putMVar m 'a'  
> 'a'
```

```
do m <- newEmptyMVar  
  forkIO (do putMVar m 'a'; putMVar m 'b')  
  takeMVar m >=> print  
  takeMVar m >=> print
```

```
do m <- newEmptyMVar  
  forkIO (do putMVar m 'a'; putMVar m 'b')  
  takeMVar m >>= print  
  takeMVar m >>= print
```

```
> do m <- newEmptyMVar; forkIO (do putMVar m 'a'; putMVar  
m 'b'); takeMVar m >>= print; takeMVar m >>= print  
'a'  
'b'
```

```
do m <- newEmptyMVar  
   takeMVar m
```

```
do m <- newEmptyMVar  
    takeMVar m
```

```
> do m <- newEmptyMVar; takeMVar m  
^CInterrupted.  
>
```

```
do m <- newEmptyMVar  
    takeMVar m
```

```
> do m <- newEmptyMVar; takeMVar m  
^CInterrupted.  
>
```

```
$ ./deadlock  
deadlock: thread blocked indefinitely in an MVar operation  
$
```

- Sometimes GHC can detect that a thread is deadlocked and send it an exception (**BlockedIndefinitelyOnMVar**)
- You can catch this exception if you want.
 - default behaviour is for the thread to die silently when it receives this exception => thread has been GC'd

```
do m <- newEmptyMVar  
  forkIO (putMVar m 'a')  
  forkIO (putMVar m 'b')  
  takeMVar m
```



```
do m <- newEmptyMVar  
  forkIO (putMVar m 'a')  
  forkIO (putMVar m 'b')  
  takeMVar m
```

```
> do m <- newEmptyMVar; forkIO (putMVar m 'a'); forkIO  
(putMVar m 'b'); takeMVar m  
'a'
```

```
do m <- newEmptyMVar  
  forkIO (putMVar m 'a')  
  forkIO (putMVar m 'b')  
  takeMVar m
```

```
> do m <- newEmptyMVar; forkIO (putMVar m 'a'); forkIO  
(putMVar m 'b'); takeMVar m  
'a'
```

- What about the thread that lost the race?
- It deadlocks, receives the **BlockedIndefinitelyOnMVar** exception, and dies.

Example: overlapping I/O

- One common use for concurrency is to overlap multiple I/O operations
 - overlapping I/O reduces latencies, and allows better use of resources



sequential I/O



overlapped I/O

- overlapping I/O is easy with threads: just do each I/O operation in a separate thread
 - the runtime takes care of making this efficient, even when there are 100k+ blocked I/O operations
- e.g. downloading multiple web pages

Downloading URLs

```
getURL :: String -> IO String
```

```
do
```

```
  m1 <- newEmptyMVar
```

```
  m2 <- newEmptyMVar
```

```
  forkIO $ do
```

```
    r <- getURL "http://www.wikipedia.org/wiki/Shovel"
```

```
    putMVar m1 r
```

```
  forkIO $ do
```

```
    r <- getURL "http://www.wikipedia.org/wiki/Spade"
```

```
    putMVar m2 r
```

```
  r1 <- takeMVar m1
```

```
  r2 <- takeMVar m2
```

```
  return (r1,r2)
```

Abstract the common pattern

- Fork a new thread to execute an IO action, and later wait for the result

```
newtype Async a = Async (MVar a)
```

```
async :: IO a -> IO (Async a)
async io = do
  m <- newEmptyMVar
  forkIO $ do r <- io; putMVar m r
  return (Async m)
```

```
wait :: Async a -> IO a
```

```
wait (Async m) = readMVar m
```

```
readMVar :: MVar a -> IO a
readMVar m = do
  a <- takeMVar m
  putMVar m a
  return a
```

Using Async....

```
do
```

```
  a1 <- async $ getURL "http://www.wikipedia.org/wiki/Shovel"  
  a2 <- async $ getURL "http://www.wikipedia.org/wiki/Spade"  
  r1 <- wait m1  
  r2 <- wait m2  
  return (r1,r2)
```

A driver to download many URLs

```
sites = ["http://www.bing.com",  
         "http://www.google.com",  
         ... ]  
  
main = mapM (async.http) sites >=> mapM wait  
  where  
    http url = do  
      (page, time) <- timeit $ getURL url  
      printf "downloaded: %s (%d bytes, %.2fs)\n"  
        url (B.length page) time
```

```
downloaded: http://www.google.com (14524 bytes, 0.17s)  
downloaded: http://www.bing.com (24740 bytes, 0.18s)  
downloaded: http://www.wikipedia.com/wiki/Spade (62586 bytes, 0.60s)  
downloaded: http://www.wikipedia.com/wiki/Shovel (68897 bytes, 0.60s)  
downloaded: http://www.yahoo.com (153065 bytes, 1.11s)
```

An MVar is many things...

- *a lock*
 - **MVar** () behaves like a lock: full is unlocked, empty is locked
 - Can be used as a mutex to protect some other shared state, or a critical region
- *one-place channel*
 - Since an **MVar** holds at most one value, it behaves like an asynchronous channel with a buffer size of one
- *a container for shared state*
 - e.g. **MVar** (Map key value)
 - convert persistent data structure into ephemeral
 - efficient (but there are other choices besides **MVar**)
- *building block*
 - **MVar** can be used to build many different concurrent data structures/abstractions...

MVar = container for shared state

```
import Data.Map as Map
import Control.Concurrent

type Name = String
type PhoneNum = String
type PhoneBook = Map Name PhoneNum

insertName :: MVar PhoneBook -> Name -> PhoneNum -> IO ()
insertName m name num = do
    book <- takeMVar m
    putMVar m (Map.insert name num book)

lookupNumber :: MVar PhoneBook -> Name -> IO (Maybe PhoneNum)
lookupNumber m name = do
    book <- readMVar m
    return (Map.lookup name book)
```

MVar = container for shared state

- taking the **MVar** locks the state
- putting the **MVar** updates the state and unlocks it again
- we can make *any* Haskell data structure into shared state this way. No need for special concurrent versions of data structures.
- Lazy evaluation can be used to avoid locking the state for too long:

```
book <- takeMVar m  
putMVar m (Map.insert name num book)
```

- Note that the insert is lazy, the **MVar** was only locked briefly (but beware of space leaks)

Unbounded buffered channels

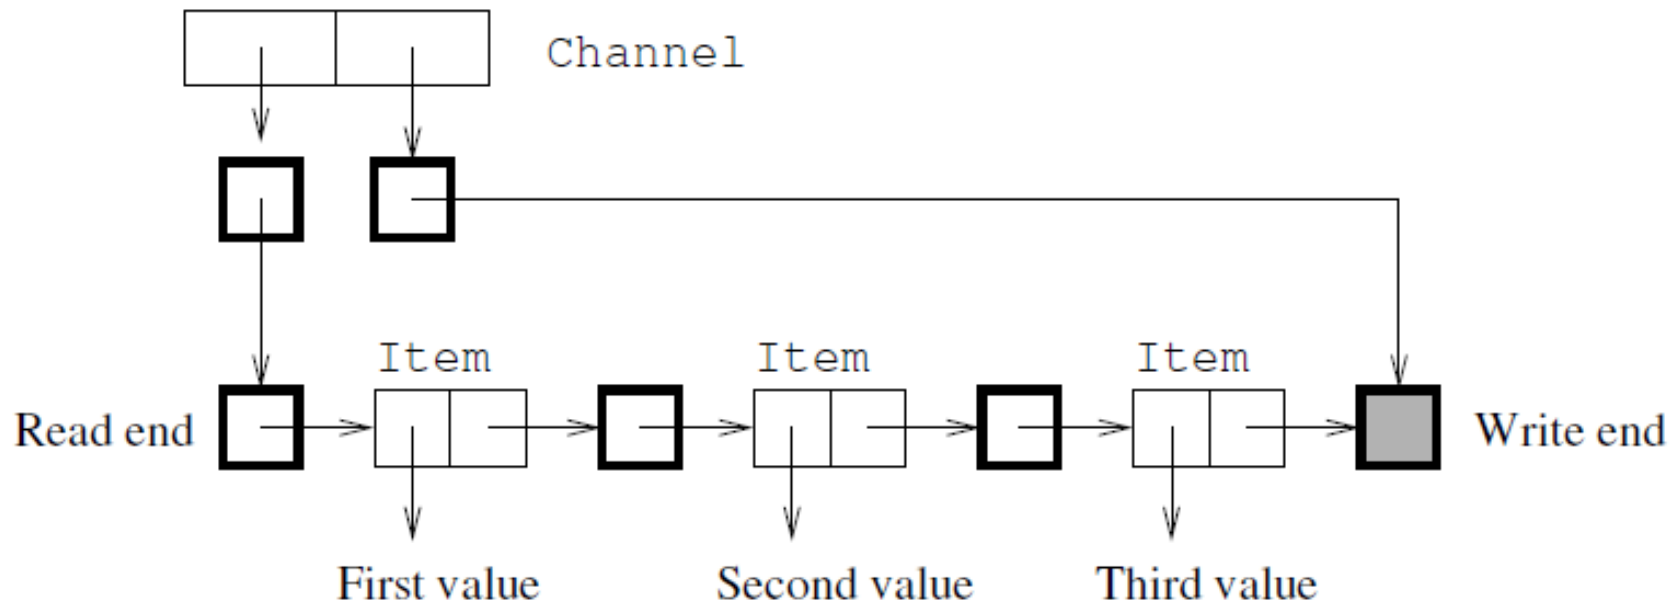
- Interface:

```
data Chan a -- abstract

newChan      :: IO (Chan a)
writeChan    :: Chan a -> a -> IO ()
readChan     :: Chan a -> IO a
```

- write does not block (indefinitely)
- we are going to implement this with MVars
- can we just use `data Chan a = MVar [a]`
- No: think about how **readChan** will block when the channel is empty
- but in both of these, writers and readers will conflict with each other


Structure of a channel



```
type Stream a = MVar (Item a)
data Item a   = Item a (Stream a)



data Chan a = Chan (MVar (Stream a))
                  (MVar (Stream a))
```

Implementation



```
newChan :: IO (Chan a)
newChan = do
  hole <- newEmptyMVar
  readVar <- newMVar hole
  writeVar <- newMVar hole
  return (Chan readVar writeVar)
```

```
writeChan :: Chan a -> a -> IO ()
writeChan (Chan _ writeVar) val = do
  new_hole <- newEmptyMVar
  old_hole <- takeMVar writeVar
  putMVar writeVar new_hole
  putMVar old_hole (Item val new_hole)
```



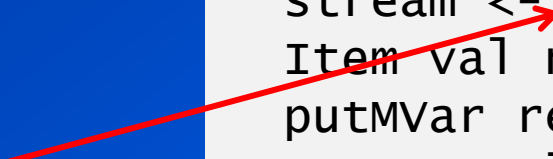
```
readChan :: Chan a -> IO a
readChan (Chan readVar _) = do
  stream <- takeMVar readVar
  Item val new <- takeMVar stream
  putMVar readVar new
  return val
```

Concurrent behaviour

- Multiple readers:

- 2nd and subsequent readers block here


```
readChan :: Chan a -> IO a
readChan (Chan readVar _) = do
    stream <- takeMVar readVar
    Item val new <- takeMVar stream
    putMVar readVar new
    return val
```



- Multiple writers:

- 2nd and subsequent writers block here

```
writeChan :: Chan a -> a -> IO ()
writeChan (Chan _ writeVar) val = do
    new_hole <- newEmptyMVar
    old_hole <- takeMVar writeVar
    putMVar writeVar new_hole
    putMVar old_hole (Item val new_hole)
```




- a concurrent read might block on **old_hole** until **writeChan** fills it in at the end

Adding more operations

- Add an operation for pushing a value onto the read end:
`unGetChan :: Chan a -> a -> IO ()`
- Doesn't seem too hard:

```
unGetChan :: Chan a -> a -> IO ()
unGetChan (Chan readVar _) val = do
  new_read_end <- newEmptyMVar
  read_end <- takeMVar readVar
  putMVar new_read_end (Item val read_end)
  putMVar readVar new_read_end
```



But...

- This doesn't work as we might expect:

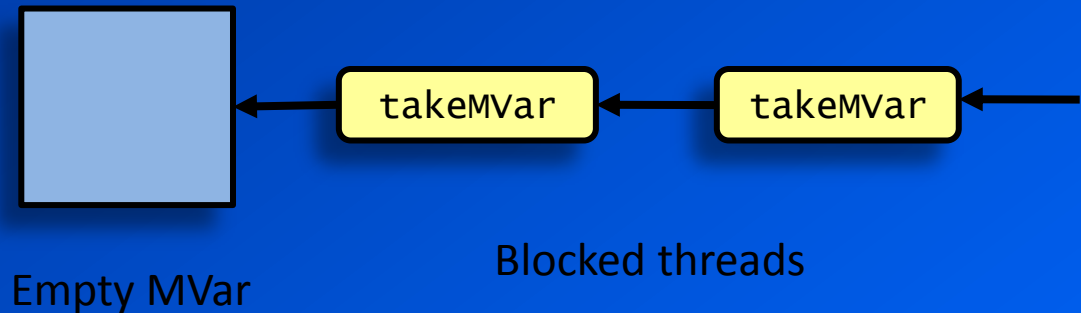
```
> c <- newChan :: IO (Chan String)
> forkIO $ do v <- readChan c; print v
ThreadId 217
> writeChan c "hello"
"hello"
> forkIO $ do v <- readChan c; print v
ThreadId 243
> unGetChan c "hello"
... blocks ....
```

- we don't expect **unGetChan** to block
- but it needs to call **takeMVar** on the read end, and the other thread is currently holding that **MVar**
- No way to fix this...
- Building larger abstractions from **MVars** can be tricky
- Software Transactional Memory is much easier (later...)

Why go to all this trouble?


- Couldn't we just build channels into the language, like Erlang, Go etc.?
 - **MVar** is a more fundamental primitive
 - You can build more than just channels with **MVar**, e.g. shared state
 - Haskell's approach is to provide simple but powerful abstractions
 - Channels *are* provided (**Control.Concurrent.Chan**)

A note about fairness




- Threads blocked on an **MVar** are processed in FIFO order
- No thread can be blocked indefinitely, provided there is a regular supply of **putMVars** (*fairness*)
- Each **putMVar** wakes exactly *one* thread, *and* performs the blocked operation atomically (*single-wakeup*)

A note about fairness



- Threads blocked on an **MVar** are processed in FIFO order
- No thread can be blocked indefinitely, provided there is a regular supply of **putMVars** (*fairness*)
- Each **putMVar** wakes exactly *one* thread, *and* performs the blocked operation atomically (*single-wakeup*)

A note about fairness




- Threads blocked on an **MVar** are processed in FIFO order
- No thread can be blocked indefinitely, provided there is a regular supply of **putMVars** (*fairness*)
- Each **putMVar** wakes exactly *one* thread, *and* performs the blocked operation atomically (*single-wakeup*)

A note about fairness



- Threads blocked on an **MVar** are processed in FIFO order
- No thread can be blocked indefinitely, provided there is a regular supply of **putMVars** (*fairness*)
- Each **putMVar** wakes exactly *one* thread, *and* performs the blocked operation atomically (*single-wakeup*)

A note about fairness



- Threads blocked on an **MVar** are processed in FIFO order
- No thread can be blocked indefinitely, provided there is a regular supply of **putMVars** (*fairness*)
- Each **putMVar** wakes exactly *one* thread, *and* performs the blocked operation atomically (*single-wakeup*)

MVars and contention

```
$ ghc fork.hs
[1 of 1] Compiling Main                ( fork.hs, fork.o )
Linking fork ...
$ ./fork | tail -c 300
AAAAAAAAABABABABABABABABABABABABABABABABABABABABABABAB
ABABABABABABABABABABABABABABABABABABABABABABABABABAB
ABABABABABABABABABABABABABABABABABABABABABABABABABAB
ABABABABABABABABABABABABABABABABABABABABABABABABABAB
```

- Fairness can lead to alternation when two threads compete for an MVar
 - thread A: takeMVar (succeeds)
 - thread B: takeMVar (blocks)
 - thread A: putMVar (succeeds, and wakes thread B)
 - thread A: takeMVar again (blocks)
 - cannot break the cycle, unless a thread is pre-empted while the MVar is full
- MVar contention can be expensive!

break...

Cancellation/interruption

(asynchronous exceptions)

Motivation

- Often we want to interrupt a thread. e.g.
 - in a web browser, the user presses “stop”
 - in a server application, we set a time-out on each client, close the connection if the client does not respond within the required time
 - if we are computing based on some input data, and the user changes the inputs via the GUI

Isn't interrupting a thread dangerous?

Isn't interrupting a thread dangerous?

- Most languages take the polling approach:
 - you have to explicitly check for interruption
 - maybe built-in support in e.g. I/O operations

Isn't interrupting a thread dangerous?

- Most languages take the polling approach:
 - you have to explicitly check for interruption
 - maybe built-in support in e.g. I/O operations
- Why?
 - because fully-asynchronous interruption is too hard to program with, in an imperative language.
 - Most code is modifying state, so asynchronous interruption will often leave state inconsistent.

Isn't interrupting a thread dangerous?

- Most languages take the polling approach:
 - you have to explicitly check for interruption
 - maybe built-in support in e.g. I/O operations
- Why?
 - because fully-asynchronous interruption is too hard to program with, in an imperative language.
 - Most code is modifying state, so asynchronous interruption will often leave state inconsistent.
- In Haskell, most computation is pure, so
 - completely safe to interrupt
 - furthermore, pure code *cannot* poll

Isn't interrupting a thread dangerous?

- Most languages take the polling approach:
 - you have to explicitly check for interruption
 - maybe built-in support in e.g. I/O operations
- Why?
 - because fully-asynchronous interruption is too hard to program with, in an imperative language.
 - Most code is modifying state, so asynchronous interruption will often leave state inconsistent.
- In Haskell, most computation is pure, so
 - completely safe to interrupt
 - furthermore, pure code *cannot* poll
- Hence, interruption in Haskell is asynchronous
 - more robust: don't have to remember to poll
 - but we do have to be careful with state in IO code

Interrupting a thread

```
throwTo :: Exception e => ThreadId -> e -> IO ()
```

- Throws the exception **e** in the given thread
- So interruption appears as an exception
 - this is good – we need exception handlers to clean up in the event of an error, and the same handlers will work for interruption too.
 - Code that is already well-behaved with respect to exceptions will be fine with interruption.

```
bracket (newTempFile "temp")  
      (\file -> removeFile file)  
      (\file -> ...)
```

- threads can also *catch* interruption exceptions and do something – e.g. useful for time-out

Example

- Let's extend the async API with cancellation
- So far we have:

```
newtype Async a = Async (MVar a)

async :: IO a -> IO (Async a)
async io = do
  m <- newEmptyMVar
  forkIO $ do r <- io; putMVar m r
  return (Async m)

wait :: Async a -> IO a
wait (Async m) = readMVar m
```

- we want to add: `cancel :: Async a -> IO ()`

- `cancel` is going to call `throwTo`, so it needs the `ThreadId`. Hence we need to add `ThreadId` to `Async`.

```
data Async a = Async ThreadId (MVar a)

async :: IO a -> IO (Async a)
async io = do
  m <- newEmptyMVar
  t <- forkIO $ do r <- io; putMVar m r
  return (Async t m)

cancel :: Async a -> IO ()
cancel (Async t _) = throwTo t ThreadKilled
```

- but what about `wait`? previously it had type:

```
wait :: Async a -> IO a
```

- but what should it return if the `Async` was cancelled?

- Cancellation is an exception, so wait should return the exception that was thrown.
- This also means that wait will correctly handle exceptions caused by errors.

```
data Async a = Async ThreadId (MVar (Either SomeException a))
```

```
async :: IO a -> IO (Async a)
```

```
async io = do
```

```
  m <- newEmptyMVar
```

```
  t <- forkIO (do r <- try action; putMVar m r)
```

```
  return (Async t m)
```

```
wait :: Async a -> IO (Either SomeException a)
```

```
wait (Async _ var) = takeMVar var
```

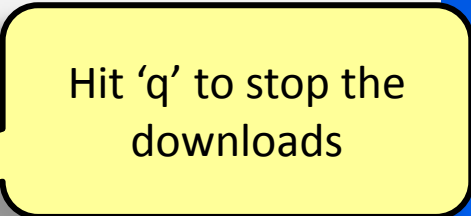
```
try :: Exception e => IO a -> IO (Either e a)
```

Example

```
main = do
  as <- mapM (async.http) sites

  forkIO $ do
    hSetBuffering stdin NoBuffering
    forever $ do
      c <- getChar
      when (c == 'q') $ mapM_ cancel as

  rs <- mapM wait as
  printf "%d/%d finished\n" (length (rights rs)) (length rs)
```



Hit 'q' to stop the downloads

```
$ ./geturlscancel  
downloaded: http://www.google.com (14538 bytes, 0.17s)  
downloaded: http://www.bing.com (24740 bytes, 0.22s)  
q2/5 finished  
$
```

- Points to note:
 - We are using a large/complicated HTTP library underneath, yet it supports interruption automatically
 - Having asynchronous interruption be the default is very powerful
 - However: dealing with truly mutable state and interruption still requires some care...

Masking asynchronous exceptions

- Problem:
 - call `takeMVar`
 - perform an operation (`f :: a -> IO a`) on the value
 - put the new value back in the `MVar`
 - if an interrupt or exception occurs anywhere, put the old value back and propagate the exception

```
catch :: Exception e => IO a -> (e -> IO a) -> IO a
```

Attempt 1

```
problem m f = do
  a <- takeMVar m
  r <- f a `catch` \e -> do putMVar m a; throw e
  putMVar m r
```

Masking asynchronous exceptions

- Problem:
 - call `takeMVar`
 - perform an operation (`f :: a -> IO a`) on the value
 - put the new value back in the `MVar`
 - if an interrupt or exception occurs anywhere, put the old value back and propagate the exception

```
catch :: Exception e => IO a -> (e -> IO a) -> IO a
```

Attempt 1

```
problem m f = do
  a <- takeMVar m
  r <- f a `catch` \e -> do putMVar m a; throw e
  putMVar m r
```

Masking asynchronous exceptions

- Problem:
 - call `takeMVar`
 - perform an operation (`f :: a -> IO a`) on the value
 - put the new value back in the `MVar`
 - if an interrupt or exception occurs anywhere, put the old value back and propagate the exception

```
catch :: Exception e => IO a -> (e -> IO a) -> IO a
```

Attempt 1

```
problem m f = do
  a <- takeMVar m
  r <- f a `catch` \e -> do putMVar m a; throw e
  putMVar m r
```


Masking asynchronous exceptions

- Problem:
 - call `takeMVar`
 - perform an operation (`f :: a -> IO a`) on the value
 - put the new value back in the `MVar`
 - if an interrupt or exception occurs anywhere, put the old value back and propagate the exception

```
catch :: Exception e => IO a -> (e -> IO a) -> IO a
```

Attempt 2

```
problem m f = do
  a <- takeMVar m
  (do r <- f a
      putMVar m r
  )
  `catch` \e -> do putMVar m a; throw e
```

Masking asynchronous exceptions

- Problem:
 - call `takeMVar`
 - perform an operation (`f :: a -> IO a`) on the value
 - put the new value back in the `MVar`
 - if an interrupt or exception occurs anywhere, put the old value back and propagate the exception

```
catch :: Exception e => IO a -> (e -> IO a) -> IO a
```

Attempt 2

```
problem m f = do
  a <- takeMVar m
  (do r <- f a
    putMVar m r
  )
  `catch` \e -> do putMVar m a; throw e
```

- Clearly we need a way to manage the delivery of asynchronous exceptions during critical sections.

- Clearly we need a way to manage the delivery of asynchronous exceptions during critical sections.
- Haskell provides `mask` for this purpose:

```
mask :: ((IO a -> IO a) -> IO b) -> IO b
```

- Clearly we need a way to manage the delivery of asynchronous exceptions during critical sections.
- Haskell provides `mask` for this purpose:

```
mask :: ((IO a -> IO a) -> IO b) -> IO b
```

- Use it like this:

```
problem m f = mask $ \restore -> do
  a <- takeMVar m
  r <- restore (f a) `catch` \e -> do putMVar m a; throw e
  putMVar m r
```

- Clearly we need a way to manage the delivery of asynchronous exceptions during critical sections.
- Haskell provides `mask` for this purpose:

```
mask :: ((IO a -> IO a) -> IO b) -> IO b
```

- Use it like this:

```
problem m f = mask $ \restore -> do  
  a <- takeMVar m  
  r <- restore (f a) `catch` \e -> do putMVar m a; throw e  
  putMVar m r
```

- `mask` takes as its argument a function `(\restore -> ...)`

- Clearly we need a way to manage the delivery of asynchronous exceptions during critical sections.
- Haskell provides `mask` for this purpose:

```
mask :: ((IO a -> IO a) -> IO b) -> IO b
```

- Use it like this:

```
problem m f = mask $ \restore -> do
  a <- takeMVar m
  r <- restore (f a) `catch` \e -> do putMVar m a; throw e
  putMVar m r
```

- `mask` takes as its argument a function `(\restore -> ...)`
- during execution of `(\restore -> ...)`, asynchronous exceptions are *masked* (blocked until the masked portion returns)

- Clearly we need a way to manage the delivery of asynchronous exceptions during critical sections.
- Haskell provides `mask` for this purpose:

```
mask :: ((IO a -> IO a) -> IO b) -> IO b
```

- Use it like this:

```
problem m f = mask $ \restore -> do  
  a <- takeMVar m  
  r <- restore (f a) `catch` \e -> do putMVar m a; throw e  
  putMVar m r
```

- `mask` takes as its argument a function `(\restore -> ...)`
- during execution of `(\restore -> ...)`, asynchronous exceptions are *masked* (blocked until the masked portion returns)
- The value passed in as the argument `restore` is a function `(:: IO a -> IO a)` that can be used to restore the previous state (unmasked or masked) inside the masked portion.

- So did this solve the problem?

```
problem m f = mask $ \restore -> do
  a <- takeMVar m
  r <- restore (f a) `catch` \e -> do putMVar m a; throw e
  putMVar m r
```

- async exceptions cannot be raised in the red portions... so we always safely put back the **MVar**, restoring the invariant
- But: what if **takeMVar** blocks?
 - We are inside **mask**, so the thread cannot be interrupted. Bad!!
 - We didn't really want to mask **takeMVar**, we only want it to atomically enter the masked state when **takeMVar** takes the value

- Solution:

- Operations that block are declared to be *interruptible*.
- An interruptible operation may receive asynchronous exceptions while blocked, even inside mask.

- How does this help?
 - **takeMVar** is now interruptible, so the thread can be interrupted while blocked
 - in general, it is now very hard to accidentally write code that is uninterruptible for long periods (it has to be in a busy loop)
- Think of **mask** as *switch to polling mode*
 - interruptible operations poll
 - we know which ops poll, so we can use exception handlers
 - asynchronous exceptions become *synchronous* inside mask

- hmm, don't we have another problem now?

```
problem m f = mask $ \restore -> do
  a <- takeMVar m
  r <- restore (f a) `catch` \e -> do putMVar m a; throw e
  putMVar m r
```

- **putMVar** is interruptible too!
- interruptible operations only receive asynchronous exceptions if they actually block
 - In this case, we can ensure that **putMVar** will never block, by requiring that all accesses to this **MVar** use a take/put pair, not just a put.
 - Alternatively, use the non-blocking version of **putMVar**, **tryPutMVar**

Async-exception safety

- IO code that uses state needs to be made safe in the presence of async exceptions
- ensure that invariants on the state are maintained if an async exception is raised.
- We make this easier by providing combinators that cover common cases.
- e.g. the function problem we saw earlier is a useful way to safely modify the contents of an **MVar**:

```
modifyMVar_ :: MVar a -> (a -> IO a) -> IO ()
```

Making Chan safe

- We had this:

```
writeChan :: Chan a -> a -> IO ()
writeChan (Chan _ writeVar) val = do
    new_hole <- newEmptyMVar
    old_hole <- takeMVar writeVar
    putMVar writeVar new_hole
    putMVar old_hole (Item val new_hole)
```

Making Chan safe

- We had this:

```
writeChan :: Chan a -> a -> IO ()  
writeChan (Chan _ writeVar) val = do  
    new_hole <- newEmptyMVar  
    old_hole <- takeMVar writeVar  
    putMVar writeVar new_hole  
    putMVar old_hole (Item val new_hole)
```

danger!

Making Chan safe

- We had this:

```
writeChan :: Chan a -> a -> IO ()
writeChan (Chan _ writeVar) val = do
  new_hole <- newEmptyMVar
  old_hole <- takeMVar writeVar
  putMVar writeVar new_hole
  putMVar old_hole (Item val new_hole)
```

danger!

- use **mask_**

```
writeChan (Chan _ writeVar) val = do
  new_hole <- newEmptyMVar
  mask_ $ do
    old_hole <- takeMVar writeVar
    putMVar writeVar new_hole
    putMVar old_hole (Item val new_hole)
```

Recap

- Basic concurrency operations:
 - forkIO
 - lightweight: make lots of them
 - MVar, takeMVar, putMVar
 - generalises mutexes, 1-place channels, state containers
- Asynchronous exceptions:
 - throwTo
 - throw an exception to another thread
 - mask
 - prevent exceptions from being thrown here for a while